

DATA SCIENCE

(Prediction of heart disease occurrence)

*Summer Internship Report Submitted in partial fulfillment
of the requirement for undergraduate degree of*

Bachelor of Technology

In

Computer Science and Engineering

By

Kannekanti Neha

221710313020

Under the Guidance of

Mr. _____

Assistant Professor



GITAM
(DEEMED TO BE UNIVERSITY)
(Estd. u/s 3 of the UGC Act, 1956)

VISAKHAPATNAM ✨ HYDERABAD ✨ BENGALURU

Department Of Computer Science And Engineering

GITAM School of Technology

GITAM (Deemed to be University)

Hyderabad-502329

June 2020

DECLARATION

I submit this industrial training work entitled “**PREDICTION OF HEART DISEASE OCCURANCE**” to GITAM (Deemed To Be University), Hyderabad in partial fulfillment of the requirements for the award of the degree of “**Bachelor of Technology**” in “**Computer Science and Engineering**”. I declare that it was carried out independently by me under the guidance of **Mr. _____**, Asst. Professor, GITAM (Deemed To Be University), Hyderabad, India.

The results embodied in this report have not been submitted to any other University or Institute for the award of any degree or diploma.

Place: HYDERABAD

Kannekanti Neha

Date:

221710313020



CERTIFICATE

This is to certify that the Industrial Training Report entitled “PREDICTION OF HEART DISEASE OCCURANCE” is being submitted by Kannekanti Neha (221710313020) in partial fulfillment of the requirement for the award of Bachelor of Technology in Computer Science and Engineering at GITAM (Deemed To Be University), Hyderabad during the academic year 2020-21

It is faithful record work carried out by her at the Computer Science and Engineering Department, GITAM University Hyderabad Campus under my guidance and supervision.

Mr. _____

Dr. S.Phani Kumar,
Professor and HOD,
Department of CSE

certificate

ACKNOWLEDGEMENT

Apart from my effort, the success of this internship largely depends on the encouragement and guidance of many others. I take this opportunity to express my gratitude to the people who have helped me in the successful completion of this internship.

I would like to thank respected Dr. N. Siva Prasad, Pro Vice Chancellor, GITAM Hyderabad and Dr. CH. Sanjay, Principal, GITAM Hyderabad.

I would like to thank respected Prof. S. Phani Kumar, Head of the Department of Computer Science Engineering for giving me such a wonderful opportunity to expand my knowledge for my own branch and giving me guidelines to present the internship report. It helped me a lot to realize what we study for.

I would like to thank the respected faculties _____ who helped me to make this internship a successful accomplishment.

I would also like to thank my friends who helped me to make my work more organized and well-structured till the end.

Kannekanti Neha
221710313020

ABSTRACT

Machine learning algorithms are used to predict the values from the data set by splitting the data set in to train and test and building Machine learning algorithms models of higher accuracy to predict the values is the primary task to be performed on heart data set My perception of understanding the given data set has been in the view of undertaking a client's requirement of overcoming the predictions on whether a person is suffering from Heart Disease or not.

Health care field has a vast amount of data, for processing those data certain techniques are used. Data science is one of the techniques often used. Heart disease is the Leading cause of death worldwide. This System predicts the arising possibilities of Heart Disease. The outcomes of this system provide the chances of occurring heart disease in terms of percentage. The datasets used are classified in terms of medical parameters. This system evaluates those parameters using data science classification technique. The datasets are processed in python programming using three main Machine Learning Algorithm namely logistic regression , knn Algorithm and Naive Bayes Algorithm which shows the best algorithm among these three in terms of accuracy level of heart disease.

Table of Contents:

CHAPTER 1: INFORMATION ABOUT DATA SCIENCE

| | |
|-----------------------------------|----|
| 1.1.1 What is Data Science? | 9 |
| 1.1.2 Need of Data Science | 10 |
| 1.1.3 Uses of Data Science | 11 |

CHAPTER 2.INFORMATION ABOUT MACHINE LEARNING

| | |
|--|----|
| 2.1.1 Introduction..... | 15 |
| 2.1.2 Importance of Machine Learning | 15 |
| 2.1.3 Uses of Machine Learning | 16 |
| 2.1.4 Types of Machine Learning | 17 |
| 2.1.5 Relation between Data Mining, Machine Learning and Deep Learning | 20 |

CHAPTER 3. INFORMATION ABOUT PYTHON

| | |
|------------------------------|----|
| 3.1 Introduction | 20 |
| 3.2 Features of Python | 21 |
| 3.3 Setup of Python | 22 |
| 3.4 Variable Types | 24 |
| 3.5 Functions | 28 |
| 3.6 OOPs Concepts | 28 |

CHAPTER 4. PROJECT: Prediction Of News as Real or Fake

| | |
|----------------------------------|----|
| 4.1 Project Requirements | 30 |
| 4.1.1 Packages Used | 30 |
| 4.1.2 Algorithms used | 31 |
| 4.2 Problem Statement | 31 |
| 4.3 Dataset Description | 32 |
| 4.4 Objective of case study..... | 33 |

CHAPTER 5. DATA PREPROCESSING

| | |
|--|----|
| 5.1 Reading the dataset | 33 |
| 5.2 Handling missing values and duplicate values..... | 34 |
| 5.3 Checking whether there are balanced no of 0's and 1's..... | 35 |

CHAPTER 6. FEATURE SELECTION

| | |
|-------------------------------|----|
| 6.1 Correlation matrix: | 36 |
| 6.2 Histogram: | 37 |

CHAPTER 7. MODEL BUILDING AND EVALUATION

| | |
|--|----|
| 7.1 Naive Bayes: | 36 |
| - 7.1.1 Brief about the algorithms used: | 36 |
| - 7.1.2 Train the Models: | 37 |
| - 7.1.3 Predicting on train data: | 37 |
| - 7.1.4 Predicting on test data: | 38 |
| 7.2 Logistic Regression:: | 39 |

| | |
|--|-----------|
| - 7.2.1 Brief about the algorithms used: ----- | 39 |
| - 7.2.2 Train the Models: ----- | 40 |
| - 7.2.3 Predicting on train data: ----- | 40 |
| - 7.2.4 Predicting on test data: ----- | 42 |
| 7.3 KNN classifier : | |
| - 7.3.1 Brief about the algorithms used: ----- | 44 |
| - 7.3.2 Train the Models: ----- | 44 |
| - 7.3.3 Predicting on train data: ----- | 49 |
| - 7.3.4 Predicting on test data: ----- | 52 |
| 8 CONCLUSION: ----- | 54 |
| 9 REFERENCES: ----- | 55 |

CHAPTER 1

INFORMATION ABOUT DATA SCIENCE

1.1.1 WHAT IS DATA SCIENCE?

Data Science is a blend of various tools, algorithms, and machine learning principles with the goal to discover hidden patterns from the raw data. How is this different from what statisticians have been doing for years?

The answer lies in the difference between explaining and predicting.

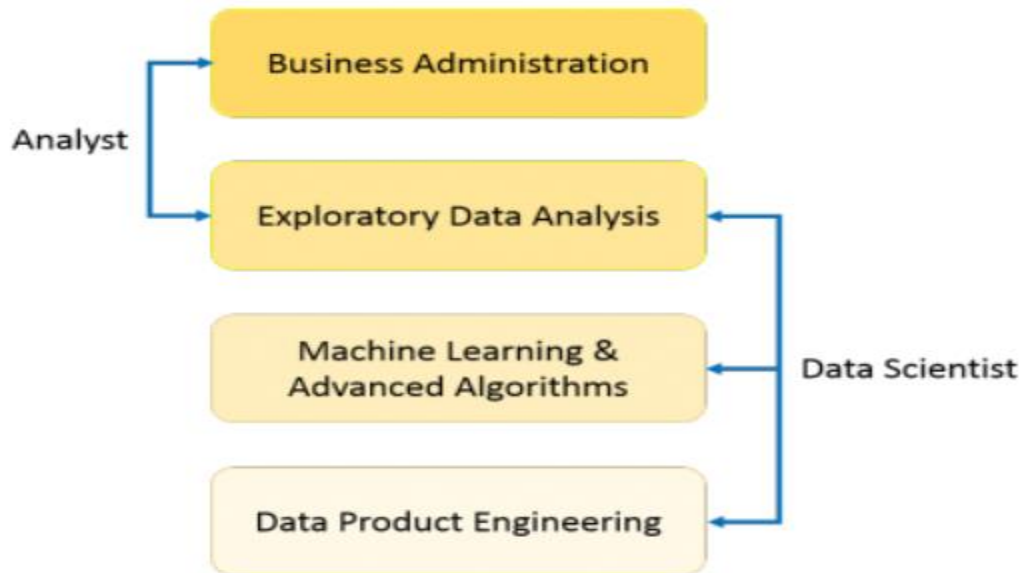


Fig.1.1 difference between explaining and predicting

As you can see from the above image, a Data Analyst usually explains what is going on by processing history of the data. On the other hand, Data Scientist not only does the exploratory analysis to discover insights from it, but also uses various advanced machine learning algorithms to identify the occurrence of a particular event in the future. A Data Scientist will look at the data from many angles, sometimes angles not known earlier.

1.1.2 NEED OF DATA SCIENCE

Data Science is primarily used to make decisions and predictions making use of predictive causal analytics, prescriptive analytics (predictive plus decision science) and machine learning.

- Predictive causal analytics – If you want a model which can predict the possibilities of a particular event in the future, you need to apply predictive causal analytics. Say, if you are providing money on credit, then the probability of customers making future credit payments on time is a matter of concern for you. Here, you can build a model which can perform predictive analytics on the payment history of the customer to predict if the future payments will be on time or not.

- Prescriptive analytics: If you want a model which has the intelligence of taking its own decisions and the ability to modify it with dynamic parameters, you certainly need prescriptive analytics for it. This relatively new field is all about providing advice. In other terms, it not only predicts but suggests a range of prescribed actions and associated outcomes.

The best example for this is Google's self-driving car which I had discussed earlier too. The data gathered by vehicles can be used to train self-driving cars. You can run algorithms on this data to bring intelligence to it. This will enable your car to take decisions like when to turn, which path to take, when to slow down or speed up.

- Machine learning for making predictions — If you have transactional data of a finance company and need to build a model to determine the future trend, then machine learning algorithms are the best bet. This falls under the paradigm of supervised learning. It is called supervised because you already have the data based on which you can train your machines. For example, a fraud detection model can be trained using a historical record of fraudulent purchases.

- Machine learning for pattern discovery — If you don't have the parameters based on which you can make predictions, then you need to find out the hidden patterns within the dataset to be able to make meaningful predictions. This is nothing but the unsupervised model as you don't have any predefined labels for grouping. The most common algorithm used for pattern discovery is Clustering.

Let's say you are working in a telephone company and you need to establish a network by putting towers in a region. Then, you can use the clustering technique to find those tower locations which will ensure that all the users receive optimum signal strength.

Let's see how the proportion of above-described approaches differ for Data Analysis as well as Data Science. As you can see in the image below, Data Analysis includes descriptive analytics and

prediction to a certain extent. On the other hand, Data Science is more about Predictive Causal Analytics and Machine Learning.

1.1.3 USES OF DATA SCIENCE

1. Medical Image Analysis

Procedures such as detecting tumors, artery stenosis, organ delineation employ various different methods and frameworks like MapReduce to find optimal parameters for tasks like lung texture classification. It applies machine learning methods, support vector machines (SVM), content-based medical image indexing, and wavelet analysis for solid texture classification.

2. Genetics & Genomics

Data Science applications also enable an advanced level of treatment personalization through research in genetics and genomics. The goal is to understand the impact of the DNA on our health and find individual biological connections between genetics, diseases, and drug response. Data science techniques allow integration of different kinds of data with genomic data in the disease research, which provides a deeper understanding of genetic issues in reactions to particular drugs and diseases. As soon as we acquire reliable personal genome data, we will achieve a deeper understanding of the human DNA. The advanced genetic risk prediction will be a major step towards more individual care.

3. Drug Development

The drug discovery process is highly complicated and involves many disciplines. The greatest ideas are often bounded by billions of testing, huge financial and time expenditure. On average, it takes twelve years to make an official submission.

Data science applications and machine learning algorithms simplify and shorten this process, adding a perspective to each step from the initial screening of drug compounds to the prediction of the success rate based on the biological factors. Such algorithms can forecast how the compound will act in the body using advanced mathematical modeling and simulations instead of the “lab

experiments”. The idea behind the computational drug discovery is to create computer model simulations as a biologically relevant network simplifying the prediction of future outcomes with high accuracy.

4. Virtual assistance for patients and customer support

Optimization of the clinical process builds upon the concept that for many cases it is not actually necessary for patients to visit doctors in person. A mobile application can give a more effective solution by *bringing the doctor to the patient instead*. The AI-powered mobile apps can provide basic healthcare support, usually as chatbots. You simply describe your symptoms, or ask questions, and then receive key information about your medical condition derived from a wide network linking symptoms to causes. Apps can remind you to take your medicine on time, and if necessary, assign an appointment with a doctor. This approach promotes a healthy lifestyle by encouraging patients to make healthy decisions, saves their time waiting in line for an appointment, and allows doctors to focus on more critical cases.

5. Internet Search

Now, this is probably the first thing that strikes your mind when you think Data Science Applications. When we speak of search, we think ‘Google’. Right? But there are many other search engines like Yahoo, Bing, Ask, AOL, and so on. All these search engines (including Google) make use of data science algorithms to deliver the best result for our searched query in a fraction of seconds. Considering the fact that, Google processes more than 20 petabytes of data every day.

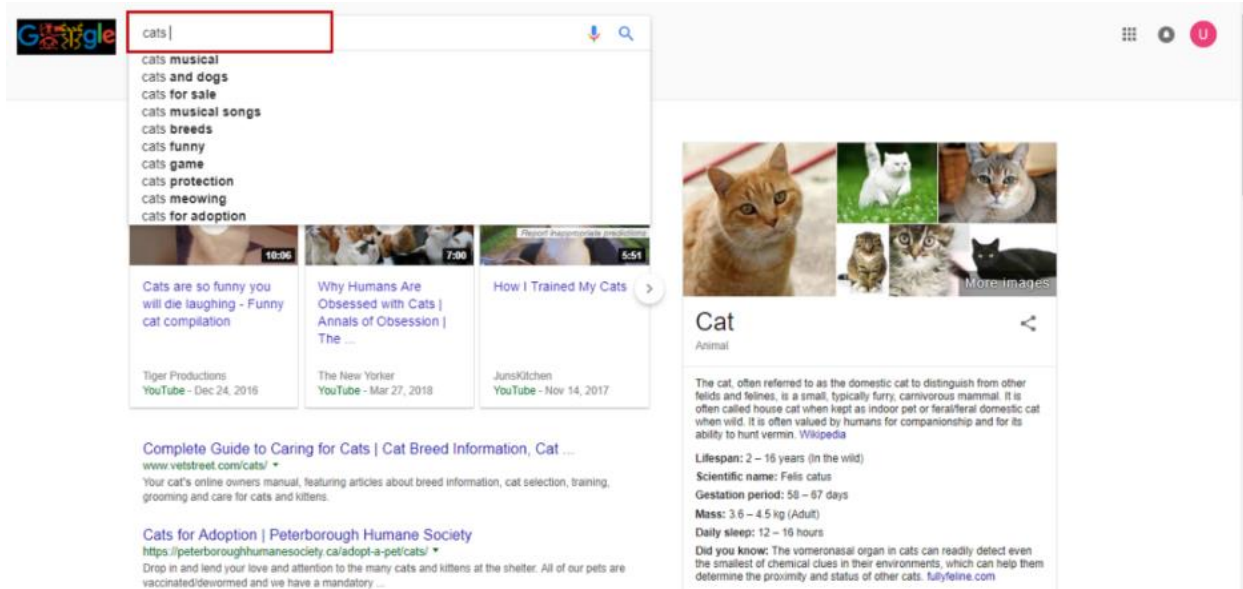


Fig.1.2 internet search

6.Targeted Advertising

If you thought Search would have been the biggest of all data science applications, here is a challenger – the entire digital marketing spectrum. Starting from the display banners on various websites to the digital billboards at the airports – almost all of them are decided by using data science algorithms.

This is the reason why digital ads have been able to get a lot higher CTR (Call-Through Rate) than traditional advertisements. They can be targeted based on a user's past behavior.

This is the reason why you might see ads of Data Science Training Programs while I see an ad of apparels in the same place at the same time.

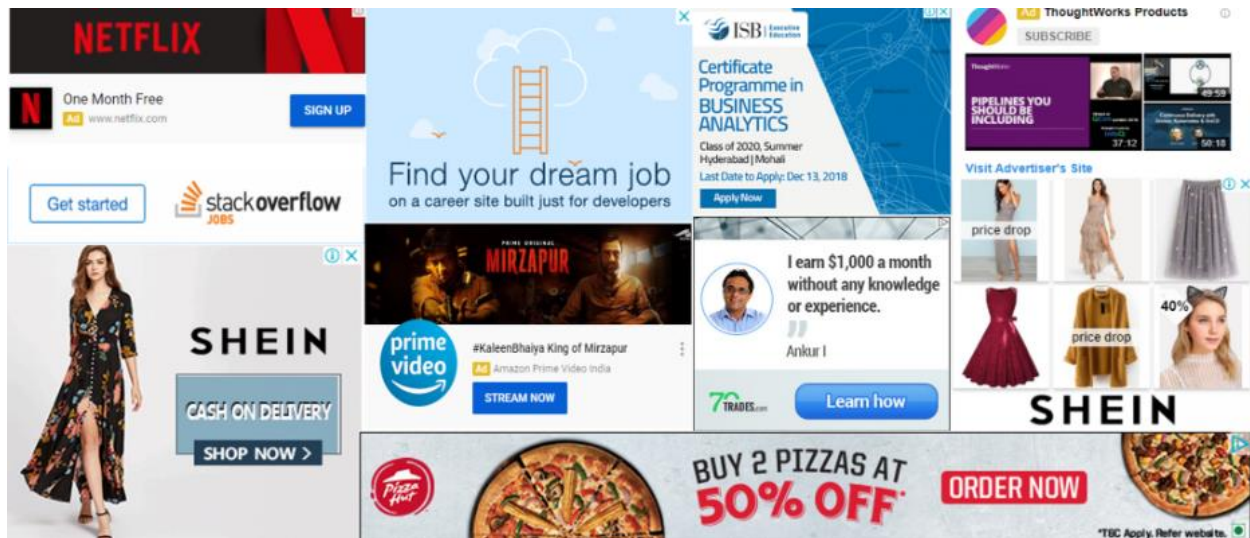


Fig 1.3 targeted advertising

CHAPTER 2

INFORMATION ABOUT MACHINE LEARNING

2.1.1 INTRODUCTION:

Machine Learning(ML) is the scientific study of algorithms and statistical models that computer systems use in order to perform a specific task effectively without using explicit instructions, relying on patterns and inference instead. It is seen as a subset of Artificial Intelligence(AI).

2.1.2 IMPORTANCE OF MACHINE LEARNING:

Consider some of the instances where machine learning is applied: the self-driving Google car, cyber fraud detection, online recommendation engines—like friend suggestions on Facebook, Netflix showcasing the movies and shows you might like, and “more items to consider” and “get yourself a little something” on Amazon—are all examples of applied machine learning. All these examples echo the vital role machine learning has begun to take in today’s data-rich world.

Machines can aid in filtering useful pieces of information that help in major advancements, and we are already seeing how this technology is being implemented in a wide variety of industries.

With the constant evolution of the field, there has been a subsequent rise in the uses, demands, and importance of machine learning. Big data has become quite a buzzword in the last few years; that’s in part due to increased sophistication of machine learning, which helps analyze those big chunks of big data. Machine learning has also changed the way data extraction, and interpretation is done by involving automatic sets of generic methods that have replaced traditional statistical techniques.

The process flow depicted here represents how machine learning works

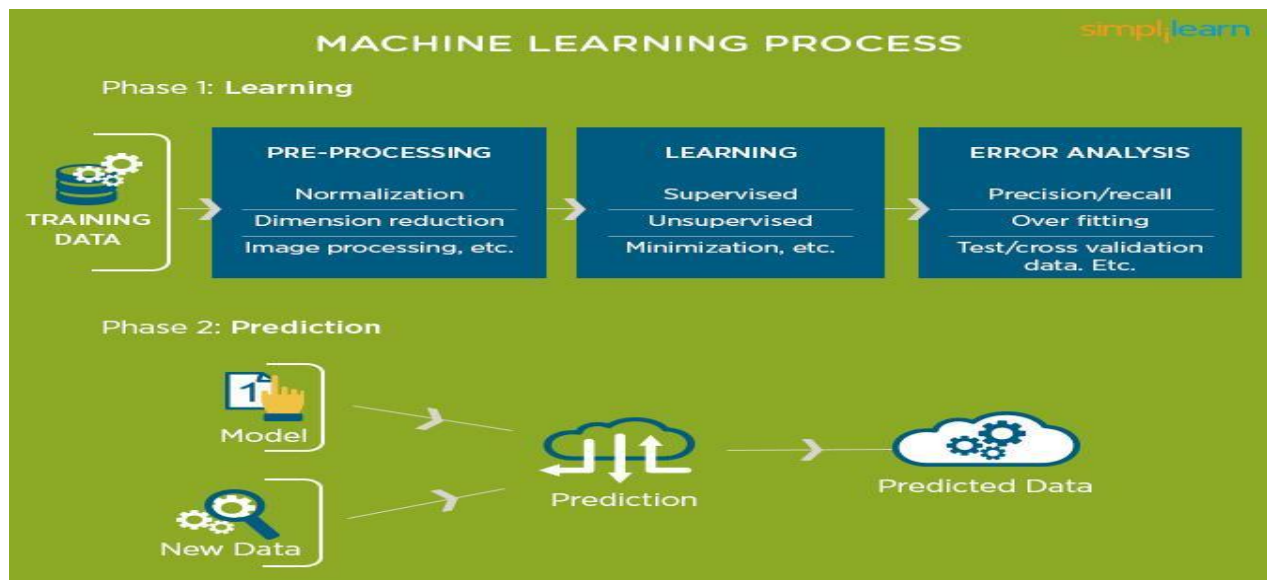


Fig 2.1 : The Process Flow

2.1.3 USES OF MACHINE LEARNING:

Earlier in this article, we mentioned some applications of machine learning. To understand the concept of machine learning better, let's consider some more examples: web search results, real-time ads on web pages and mobile devices, email spam filtering, network intrusion detection, and pattern and image recognition. All these are by-products of applying machine learning to analyze huge volumes of data. Traditionally, data analysis was always being characterized by trial and error, an approach that becomes impossible when data sets are large and heterogeneous. Machine learning comes as the solution to all this chaos by proposing clever alternatives to analyzing huge volumes of data.

By developing fast and efficient algorithms and data-driven models for real-time processing of data, machine learning can produce accurate results and analysis.

2.1.4 TYPES OF LEARNING ALGORITHMS:

The types of machine learning algorithms differ in their approach, the type of data they input and output, and the type of task or problem that they are intended to solve.

Supervised Learning :

When an algorithm learns from example data and associated target responses that can consist of numeric values or string labels, such as classes or tags, in order to later predict the correct response when posed with new examples comes under the category of supervised learning.

Supervised machine learning algorithms uncover insights, patterns, and relationships from a labelled training dataset – that is, a dataset that already contains a known value for the target variable for each record. Because you provide the machine learning algorithm with the correct answers for a problem during training, it is able to “learn” how the rest of the features relate to the target, enabling you to uncover insights and make predictions about future outcomes based on historical data.

Examples of Supervised Machine Learning Techniques are Regression, in which the algorithm returns a numerical target for each example, such as how much revenue will be generated from a new marketing campaign.

Classification, in which the algorithm attempts to label each example by choosing between two or more different classes. Choosing between two classes is called binary classification, such as determining whether or not someone will default on a loan. Choosing between more than two classes is referred to as multiclass classification.

Unsupervised Learning:

When an algorithm learns from plain examples without any associated response, leaving to the algorithm to determine the data patterns on its own. This type of algorithm tends to restructure the data into something else, such as new features that may represent a class or a new series of uncorrelated values. They are quite useful in providing

humans with insights into the meaning of data and new useful inputs to supervised machine learning algorithms.

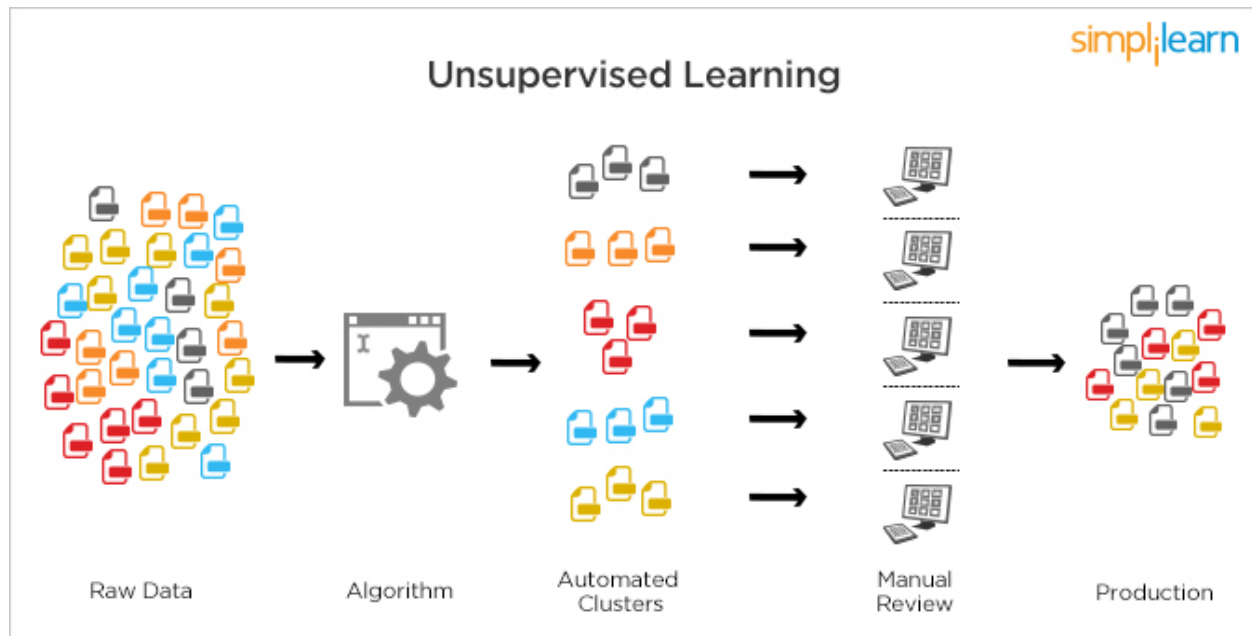


Fig 2.2 : Unsupervised Learning

Popular techniques where unsupervised learning is used also include self-organizing maps, nearest neighbor mapping, singular value decomposition, and k-means clustering. Basically, online recommendations, identification of data outliers, and segment text topics are all examples of unsupervised learning.

Semi Supervised Learning:

As the name suggests, semi-supervised learning is a bit of both supervised and unsupervised learning and uses both labeled and unlabeled data for training. In a typical scenario, the algorithm would use a small amount of labeled data with a large amount of unlabeled data.

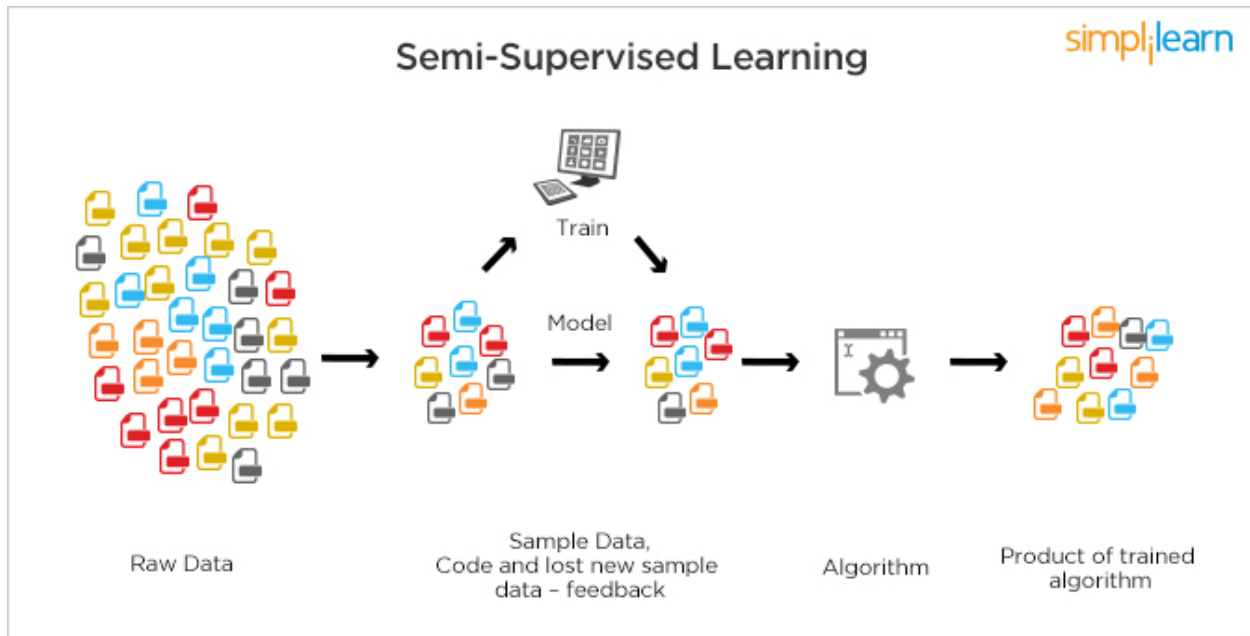


Fig 2.3 : Semi Supervised Learning

2.1.5 RELATION BETWEEN DATA MINING, MACHINE LEARNING AND DEEP LEARNING:

Machine learning and data mining use the same algorithms and techniques as data mining, except the kinds of predictions vary. While data mining discovers previously unknown patterns and knowledge, machine learning reproduces known patterns and knowledge—and further automatically applies that information to data, decision-making, and actions.

Deep learning, on the other hand, uses advanced computing power and special types of neural networks and applies them to large amounts of data to learn, understand, and identify complicated patterns. Automatic language translation and medical diagnoses are examples of deep learning.

CHAPTER 3

INFORMATION ABOUT PYTHON

Basic programming language used for machine learning is : PYTHON

3.1 INTRODUCTION TO PYHTON:

- Python is a high-level, interpreted, interactive and object-oriented scripting language.
- Python is a general purpose programming language that is often applied in scripting roles
- Python is Interpreted: Python is processed at runtime by the interpreter. You do not need to compile your program before executing it. This is like PERL and PHP.
- Python is Interactive: You can sit at a Python prompt and interact with the interpreter directly to write your programs.
- Python is Object-Oriented: Python supports the Object-Oriented style or technique of programming that encapsulates code within objects.

HISTORY OF PYTHON:

- Python was developed by GUIDO VAN ROSSUM in early 1990's
- Its latest version is 3.7 , it is generally called as python

3.2 FEATURES OF PYTHON:

- Easy-to-learn: Python has few keywords, simple structure, and a clearly defined syntax, This allows the student to pick up the language quickly.

- Easy-to-read: Python code is more clearly defined and visible to the eyes.
- Easy-to-maintain: Python's source code is fairly easy-to-maintaining.
- A broad standard library: Python's bulk of the library is very portable and cross-platform compatible on UNIX, Windows, and Macintosh.
- Portable: Python can run on a wide variety of hardware platforms and has the same interface on all platforms.
- Extendable: You can add low-level modules to the Python interpreter. These modules enable programmers to add to or customize their tools to be more efficient.
- Databases: Python provides interfaces to all major commercial databases.
- GUI Programming: Python supports GUI applications that can be created and ported to many system calls, libraries and windows systems, such as Windows MFC, Macintosh, and the X Window system of Unix.

3.3 HOW TO SETUP PYTHON:

- Python is available on a wide variety of platforms including Linux and Mac OS X. Let's understand how to set up our Python environment.
- The most up-to-date and current source code, binaries, documentation, news, etc., is available on the official website of Python.

Installation(using python IDLE):

- Installing python is generally easy, and nowadays many Linux and Mac OS distributions include a recent python.
- Download python from www.python.org

- When the download is completed, double click the file and follow the instructions to install it.
- When python is installed, a program called IDLE is also installed along with it. It provides a graphical user interface to work with python.

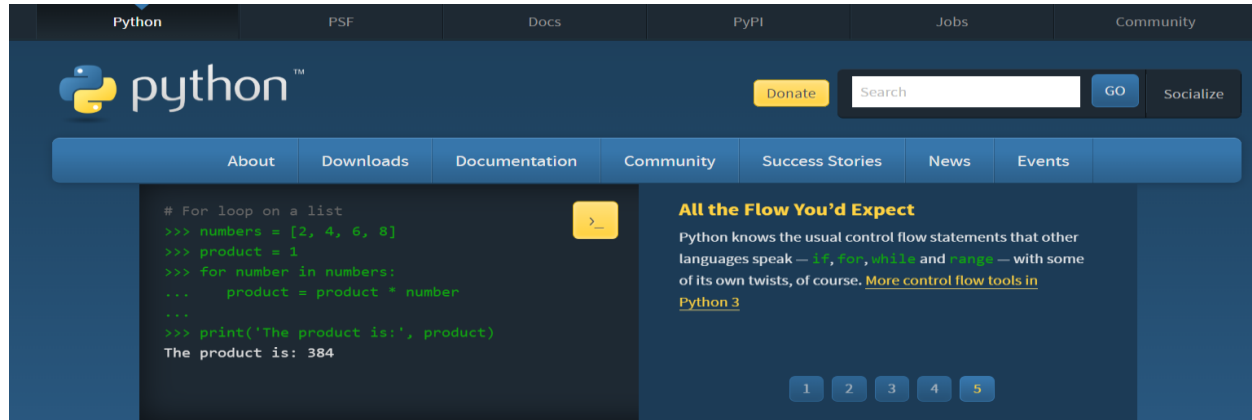


Fig 3.1 : Python download

Installation(using Anaconda):

- Python programs are also executed using Anaconda.
- Anaconda is a free open source distribution of python for large scale data processing, predictive analytics and scientific computing.
- Conda is a package manager quickly installs and manages packages.
- **In WINDOWS:**
 - In windows
 - Step 1: Open Anaconda.com/downloads in web browser.
 - Step 2: Download python 3.4 version for (32-bitgraphic installer/64 -bit graphic installer)
 - Step 3: select installation type(all users)
 - Step 4: Select path(i.e. add anaconda to path & register anaconda as default python 3.4) next click install and next click finish

- Step 5: Open jupyter notebook (it opens in default browser).

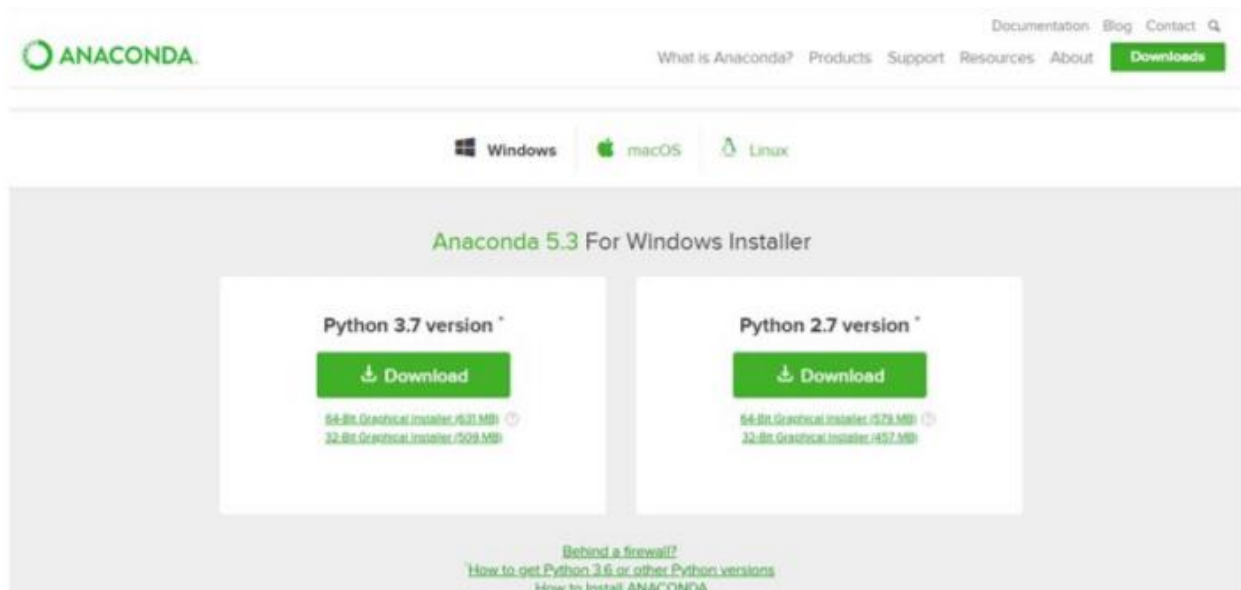


Fig 3.2 : Anaconda download




Fig 3.3 : Jupyter notebook

3.4 PYTHON VARIABLE TYPES:

- Variables are nothing but reserved memory locations to store values. This means that when you create a variable you reserve some space in memory.
- Variables are nothing but reserved memory locations to store values.
- Based on the data type of a variable, the interpreter allocates memory and decides what can be stored in the reserved memory.
- Python variables do not need explicit declaration to reserve memory space. The declaration happens automatically when you assign a value to a variable.

- Python has various standard data types that are used to define the operations possible on them and the storage method for each of them.

- Python has five standard data types –

- o Numbers

- o Strings

- o Lists

- o Tuples

- o Dictionary

Python Numbers:

- Number data types store numeric values. Number objects are created when you assign a value to them.

- Python supports four different numerical types – int (signed integers) long (long integers, they can also be represented in octal and hexadecimal) float (floating point real values) complex (complex numbers).

Python Strings:

- Strings in Python are identified as a contiguous set of characters represented in the quotation marks.

- Python allows for either pairs of single or double quotes.

- Subsets of strings can be taken using the slice operator ([] and [:]) with indexes starting at 0 in the beginning of the string and working their way from -1 at the end.

- The plus (+) sign is the string concatenation operator and the asterisk (*) is the repetition operator.

Python Lists:

- Lists are the most versatile of Python's compound data types
- A list contains items separated by commas and enclosed within square brackets.
- To some extent, lists are similar to arrays in C. One difference between them is that all the items belonging to a list can be of different data type.
- The values stored in a list can be accessed using the slice operator ([] and [:]) with indexes starting at 0 in the beginning of the list and working their way to end -1.
- The plus (+) sign is the list concatenation operator, and the asterisk (*) is the repetition operator.

Python Tuples:

- A tuple is another sequence data type that is similar to the list.
- A tuple consists of a number of values separated by commas. Unlike lists, however, tuples are enclosed within parentheses.
- The main differences between lists and tuples are: Lists are enclosed in brackets ([]) and their elements and size can be changed, while tuples are enclosed in parentheses (()) and cannot be updated.

- Tuples can be thought of as read-only lists.

● For example – Tuples are fixed size in nature whereas lists are dynamic. In other words, a tuple is immutable whereas a list is mutable. You can't add elements to a tuple. Tuples have no append or extend method. You can't remove elements from a tuple. Tuples have no remove or pop method.

Python Dictionary:

- Python's dictionaries are kind of hash table type. They work like associative arrays or hashes found in Perl and consist of key-value pairs. A dictionary key can be almost any Python type, but are usually numbers or strings. Values, on the other hand, can be any arbitrary Python object.

- Dictionaries are enclosed by curly braces ({ }) and values can be assigned and accessed using square braces ([]).

- You can use numbers to "index" into a list, meaning you can use numbers to find out what's in lists. You should know this about lists by now, but make sure you understand that you can only use numbers to get items out of a list.

- What a dict does is let you use anything, not just numbers. Yes, a dict associates one thing to another, no matter what it is.

3.5 PYTHON FUNCTION:

Defining a Function:

You can define functions to provide the required functionality. Here are simple rules to define a function in Python. Function blocks begin with the keyword `def` followed by the function name and parentheses (i.e.()).

Any input parameters or arguments should be placed within these parentheses. You can also define parameters inside these parentheses.

The code block within every function starts with a colon (:) and is indented. The statement `return [expression]` exits a function, optionally passing back an expression to the caller. A return statement with no arguments is the same as `return None`.

Calling a Function:

Defining a function only gives it a name, specifies the parameters that are to be included in the function and structures the blocks of code. Once the basic structure of a function is finalized, you can execute it by calling it from another function or directly from the Python prompt.

3.6 PYTHON USING OOP's CONCEPTS:

Class:


- **Class:** A user-defined prototype for an object that defines a set of attributes that characterize any object of the class. The attributes are data members (class variables and instance variables) and methods, accessed via dot notation.
- **Class variable:** A variable that is shared by all instances of a class. Class variables are defined within a class but outside any of the class's methods. Class variables are not used as frequently as instance variables are.

- Data member: A class variable or instance variable that holds data associated with a class and its objects.

- Instance variable: A variable that is defined inside a method and belongs only to the current instance of a class.

- Defining a Class:

- o We define a class in a very similar way how we define a function.
 - o Just like a function ,we use parentheses and a colon after the class name(i.e. ():) when we define a class. Similarly, the body of our class is indented like a functions body is.



The image shows two side-by-side code snippets. The left snippet shows a function definition: `def my_function():` followed by two indented lines of comments: `# the details of the` and `# function go here`. The right snippet shows a class definition: `class MyClass():` followed by two indented lines of comments: `# the details of the` and `# class go here`. Both snippets are presented in a light blue box with a subtle border.

Fig 3.4 : Defining a Class

`__init__` method in Class:

- The init method — also called a constructor — is a special method that runs when an instance is created so we can perform any tasks to set up the instance.

- The init method has a special name that starts and ends with two underscores: `__init__()`.

CHAPTER 4

PROJECT NAME(INFORMATION ABOUT THE PROJECT)

4.1 PROJECT REQUIREMENTS

4.1.1 Packages Used

The packages used are:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from matplotlib import rcParams
from matplotlib.cm import rainbow
%matplotlib inline
```

Fig 4.1.1 importing packages

For processing the data, I'll import a few libraries. To split the available dataset for testing and training, I'll use the `train_test_split` method. To scale the features, I am using `StandardScaler`.

CodeText

```
[ ] from sklearn.model_selection import train_test_split
    from sklearn.preprocessing import StandardScaler
```

4.1.2 Algorithms Used

- Naive Bayes Classifier
- Logistic Regression

- K Neighbors Classifier

4.2 PROBLEM STATEMENT

The problem we are going to solve here is the predictions on whether a person is suffering from Heart Disease or not. , taking into consideration the features of the dataset heart.csv

4.3 DATASET DESCRIPTION

The dataset consists of the following features:

#display the first few rows of dataset

```
[ ] heart.head()
```

| | age | sex | cp | trestbps | chol | fbs | restecg | thalach | exang | oldpeak | slope | ca | thal | target |
|---|-----|-----|----|----------|------|-----|---------|---------|-------|---------|-------|----|------|--------|
| 0 | 52 | 1 | 0 | 125 | 212 | 0 | 1 | 168 | 0 | 1.0 | 2 | 2 | 3 | 0 |
| 1 | 53 | 1 | 0 | 140 | 203 | 1 | 0 | 155 | 1 | 3.1 | 0 | 0 | 3 | 0 |
| 2 | 70 | 1 | 0 | 145 | 174 | 0 | 1 | 125 | 1 | 2.6 | 0 | 0 | 3 | 0 |
| 3 | 61 | 1 | 0 | 148 | 203 | 0 | 1 | 161 | 0 | 0.0 | 2 | 1 | 3 | 0 |
| 4 | 62 | 0 | 0 | 138 | 294 | 1 | 1 | 106 | 0 | 1.9 | 1 | 3 | 2 | 0 |

```
[ ] heart.tail()
```

| | age | sex | cp | trestbps | chol | fbs | restecg | thalach | exang | oldpeak | slope | ca | thal | target |
|------|-----|-----|----|----------|------|-----|---------|---------|-------|---------|-------|----|------|--------|
| 1020 | 59 | 1 | 1 | 140 | 221 | 0 | 1 | 164 | 1 | 0.0 | 2 | 0 | 2 | 1 |
| 1021 | 60 | 1 | 0 | 125 | 258 | 0 | 0 | 141 | 1 | 2.8 | 1 | 1 | 3 | 0 |
| 1022 | 47 | 1 | 0 | 110 | 275 | 0 | 0 | 118 | 1 | 1.0 | 1 | 1 | 2 | 0 |
| 1023 | 50 | 0 | 0 | 110 | 254 | 0 | 0 | 159 | 0 | 0.0 | 2 | 0 | 2 | 1 |
| 1024 | 54 | 1 | 0 | 120 | 188 | 0 | 1 | 113 | 0 | 1.4 | 1 | 1 | 3 | 0 |

Fig 4.3 dataset description

4.4 OBJECTIVE OF CASE STUDY

To get a better understanding of whether a predictions on whether a person is suffering from Heart Disease or not by considering the features of the data and provide the client with desired results.

CHAPTER 5

DATA PREPROCESSING

5.1 READING THE DATASET

Pandas in python provide an interesting method `read_csv()`. The `read_csv` function reads the entire dataset from a comma separated values file and we can assign it to a DataFrame to which all the operations can be performed. It helps us to access each and every row as well as columns and each and every value can be access using the dataframe.

```
[ ] heart = pd.read_csv('heart.csv')
```

Fig 5.1 reading data

```
[ ] heart.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1025 entries, 0 to 1024
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         1025 non-null   int64
1   sex         1025 non-null   int64
2   cp          1025 non-null   int64
3   trestbps    1025 non-null   int64
4   chol        1025 non-null   int64
5   fbs         1025 non-null   int64
6   restecg     1025 non-null   int64
7   thalach     1025 non-null   int64
8   exang       1025 non-null   int64
9   oldpeak     1025 non-null   float64
10  slope       1025 non-null   int64
11  ca          1025 non-null   int64
12  thal        1025 non-null   int64
13  target      1025 non-null   int64
dtypes: float64(1), int64(13)
memory usage: 112.2 KB
```

Looks like the dataset has a total of 303 rows and there are no missing values. There are a total of 13 features along with one target value

5.2 HANDLING MISSING VALUES AND DUPLICATE VALUES

We can find the number of missing values and duplicated values in each column using `isnull()` and `duplicated()` functions respectively. For our dataset no missing values or duplicated values were found.


```
[ ] heart.isnull().sum() # no null values in the dataset
```

```
↳ age      0
   sex      0
   cp       0
   trestbps 0
   chol     0
   fbs      0
   restecg  0
   thalach  0
   exang     0
   oldpeak  0
   slope    0
   ca       0
   thal     0
   target   0
   dtype: int64
```

```
[ ] print('duplicated entries: {}'.format(heart.duplicated().sum()))
```

```
↳ duplicated entries: 723
```

```
[ ] heart.drop_duplicates(inplace = True)
   heart.shape
```

```
↳ (302, 14)
```

```
[ ] print('duplicate entries: {}'.format(heart.duplicated().sum())) #no duplicates values in the dataset
```

```
↳ duplicate entries: 0
```

Fig 5.2 handling missing values and duplicates

5.3 checking whether there are balanced number of 0's and 1's

```
[ ] heart['target'].value_counts()
```

```
↳ 1      526
   0      499
   Name: target, dtype: int64
```

There is no balanced count of 0 and 1 so, to balance them by equal values iam using down sampling

```
[ ] # Separate input features (X) and target variable (y)
    y = heart.target
    X = heart.drop('target', axis=1)

[ ] from sklearn.utils import resample

    # Separate majority and minority classes
    heart_majority = heart[heart.target==0]
    heart_minority = heart[heart.target==1]

    # Downsample majority class
    heart_majority_downsampled = resample(heart_majority,
                                         replace=True,      # sample without replacement
                                         n_samples=526,      # to match minority class
                                         random_state=123)    # reproducible results

    # Combine minority class with downsampled majority class
    heart_downsampled = pd.concat([heart_majority_downsampled, heart_minority])

    # Display new class counts
    heart_downsampled.target.value_counts()
```

```
1    526
0    526
Name: target, dtype: int64
```

Fig 5.3 Balanced DataSet


CHAPTER 6

FEATURE SELECTION

6.1 Correlation Matrix

we can use visualizations to better understand our data and then look at any processing we might want to do.

```
rcParams['figure.figsize'] = 20, 14  
plt.matshow(heart.corr())  
plt.yticks(np.arange(heart.shape[1]), heart.columns)  
plt.xticks(np.arange(heart.shape[1]), heart.columns)  
plt.colorbar()
```



Taking a look at the correlation matrix above, it's easy to see that a few features have negative correlation with the target value while some have positive. Next, I'll take a look at the histograms for each variable.

6.2 Histogram

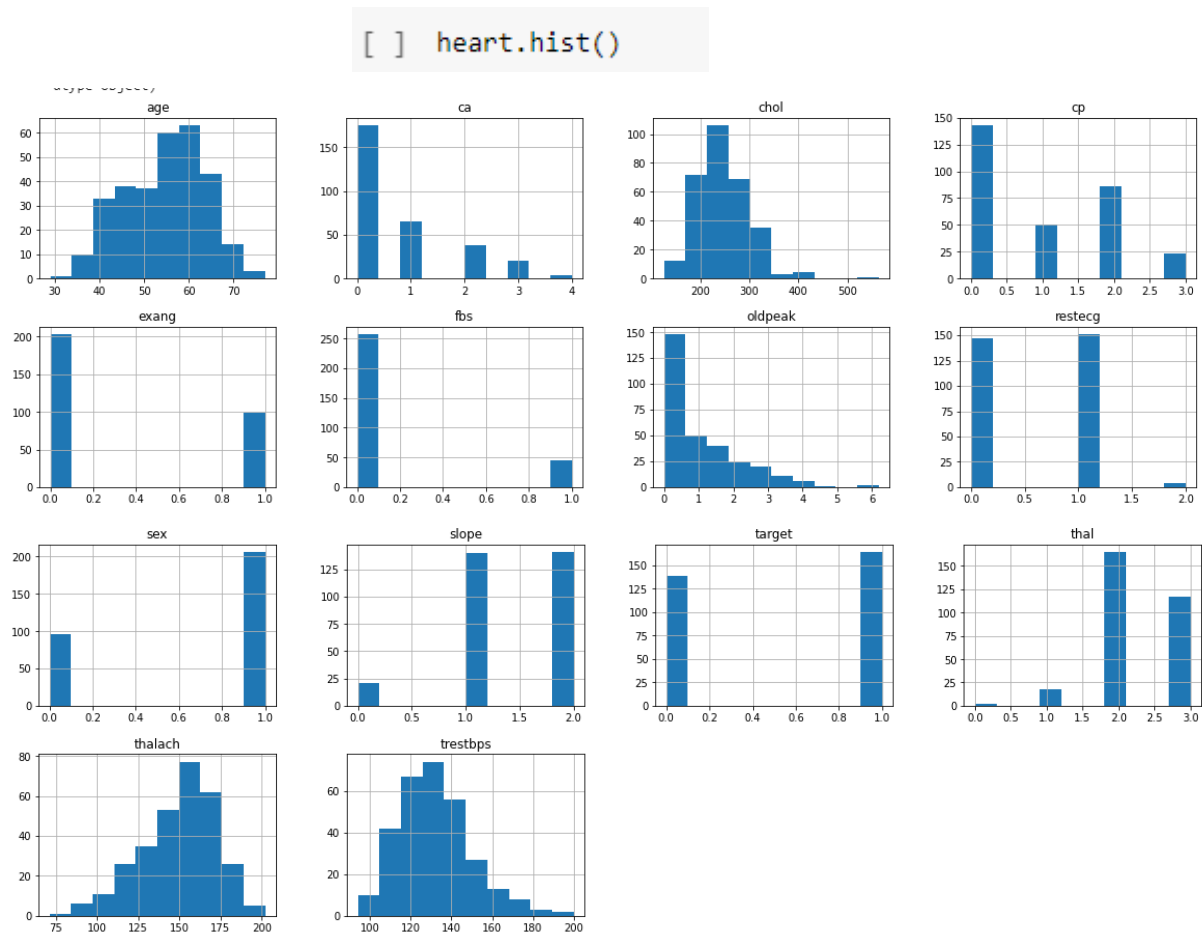


Fig 6.2. Taking a look at the histograms above, I can see that each feature has a different range of distribution. Thus, using scaling before our predictions should be of great use. Also, the categorical features do stand out.

6.3 It's always a good practice to work with a dataset where the target classes are of approximately equal size. Thus, let's check for the same.

```
[ ] rcParams['figure.figsize'] = 8,6
plt.bar(heart['target'].unique(), heart['target'].value_counts(), color = ['red', 'green'])
plt.xticks([0, 1])
plt.xlabel('Target Classes')
plt.ylabel('Count')
plt.title('Count of each Target Class')
```

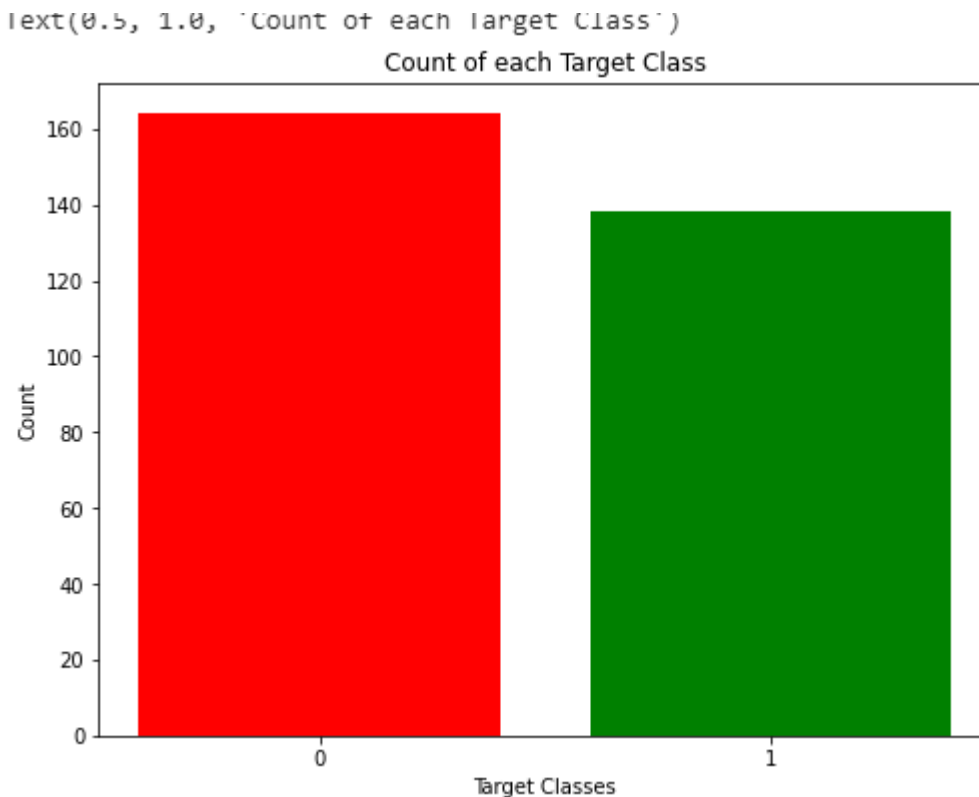


Fig 6.3 :The two classes are not exactly 50% each but the ratio is good enough to continue without dropping/increasing our data.

CHAPTER 7

MODEL BUILDING AND EVALUATION

Approach I : Naive Bayes Classifier

7.1.1 Brief about the algorithms used

In statistics, Naïve Bayes classifiers are a family of simple "probabilistic classifier" based on applying Bayes' theorem with strong (naïve) independence assumptions between the features. They are among the simplest Bayesian network models. But they could be coupled with Kernel density estimation and achieve higher accuracy levels.

Naïve Bayes classifiers are highly scalable, requiring a number of parameters linear in the number of variables (features/predictors) in a learning problem. maximum-likelihood training can be done by evaluating a closed-form expression which takes linear time, rather than by expensive iterative approximation as used for many other types of classifiers.

In the statistics and computer science literature, naive Bayes models are known under a variety of names, including simple Bayes and independence Bayes. All these names reference the use of Bayes' theorem in the classifier's decision rule, but naïve Bayes is not (necessarily) a Bayesian method.

-> Apply naive bayes algorithm to the data by importing BernoulliNB from sklearn.naive_bayes, then create an object and use the fit method on training data.

```
[ ] # Apply the naive Bayes Algorithm
    # Import BernNB
    from sklearn.naive_bayes import BernoulliNB
    # creating an object for BernNB
    model_BernNB = BernoulliNB()
```

```
[ ] # Applying the algorithm to the data
    # objectName.fit(Input,Output)
    model_BernNB.fit(X_train, y_train)
```

```
☞ BernoulliNB(alpha=1.0, binarize=0.0, class_prior=None, fit_prior=True)
```

Fig 7.1.1. applying naive bayes algorithm

7.1.2 Predicting on Train Data

We can predict the training data using predict function on X_train

```
[ ] y_train_pred = model_BernNB.predict(X_train)

[ ] # compare the actual values(y_test) with predicted values(y_test_pred)
    from sklearn.metrics import confusion_matrix,classification_report
    confusion_matrix(y_train,y_train_pred)

☞ array([[80, 12],
         [12, 98]])
```

Fig 7.1.2. predicting on train data

7.1.3 Print the classification report to know the accuracy score of training data

```
[ ] print(classification_report(y_train,y_train_pred))

☞
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.87 | 0.87 | 0.87 | 92 |
| 1 | 0.89 | 0.89 | 0.89 | 110 |
| accuracy | | | 0.88 | 202 |
| macro avg | 0.88 | 0.88 | 0.88 | 202 |
| weighted avg | 0.88 | 0.88 | 0.88 | 202 |

Fig 7.1.3. classification report for training data(CountVectorizer)

We got an accuracy score of about 88% to the training data which is considered to be a good score.

Now let's check the same for test data

7.1.4 Predicting on Test Data

We can predict the testing data using predict function on X_test

```
[ ] # Applying the algorithm to the data
    # objectName.fit(Input,Output)
    model_BernNB.fit(X_test, y_test)

↳ BernoulliNB(alpha=1.0, binarize=0.0, class_prior=None, fit_prior=True)

[ ] y_test_pred = model_BernNB.predict(X_test)

[ ] # compare the actual values(y_test) with predicted values(y_test_pred)
    from sklearn.metrics import confusion_matrix,classification_report
    confusion_matrix(y_test,y_test_pred)

↳ array([[37,  9],
        [10, 44]])
```

Fig 7.1.4. predicting on test data

7.1.5 Print the classification report to know the accuracy score of testing data

```
[ ] print(classification_report(y_test,y_test_pred))

↳
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.79 | 0.80 | 0.80 | 46 |
| 1 | 0.83 | 0.81 | 0.82 | 54 |
| accuracy | | | 0.81 | 100 |
| macro avg | 0.81 | 0.81 | 0.81 | 100 |
| weighted avg | 0.81 | 0.81 | 0.81 | 100 |

Fig 7.1.5 classification report for testing data

We got an accuracy score of around 81%, therefore we can say that it is a best fit and the model predicted very well.

Approach II : Logistic Regression

7.2.1 Brief about the algorithms used

In statistics, the **logistic model** (or **logit model**) is used to model the probability of a certain class or event existing such as pass/fail, win/lose, alive/dead or healthy/sick. This can be extended to model several classes of events such as determining whether an image contains a cat, dog, lion, etc. Each object being detected in the image would be assigned a probability between 0 and 1, with a sum of one.

7.2.2 Train the model

Import LogisticRegression from sklearn.linear_model and create an object. Fit the model on training data

```
[ ] from sklearn.linear_model import LogisticRegression
    log_reg = LogisticRegression() # creating an object for Logistic Regression

[ ] ## We have to apply this object(log_reg) to the training data
    log_reg.fit(X_train, y_train) # with help of fit method we are fitting the
                                ##Logistic Regression on training data
    ## objectName.fit(InputData, OutputData)

[ ] LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
    intercept_scaling=1, l1_ratio=None, max_iter=100,
    multi_class='auto', n_jobs=None, penalty='l2',
    random_state=None, solver='lbfgs', tol=0.0001, verbose=0,
    warm_start=False)
```

Fig 7.2.2. importing logistic regression

7.2.3 Predicting on Train data

We can predict the training data using predict function on X_train

```
[ ] y_train_pred = log_reg.predict(X_train)
```

```
[ ] y_train == y_train_pred
```

```
426    True
58     False
363    True
19     True
299    True
...
425    True
271    True
143    True
50     True
232    True
Name: target, Length: 202, dtype: bool
```

Fig 7.2.3. predicting on train data

Import confusion_matrix and classification_report from sklearn.metrics

Visualizing the confusion matrix using heat map

```
[ ] from sklearn.metrics import confusion_matrix, accuracy_score
conf = confusion_matrix(y_train, y_train_pred)
conf
sns.heatmap(confusion_matrix(y_train, y_train_pred), annot=True, fmt='3.0f', annot_kws={'size': '10', "ha": "right", "va": "baseline"})
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f88d9ba3ef0>
```




Fig 7.2.3.b visualizing confusion matrix of train data

Find the accuracy score

```
[ ] from sklearn.metrics import accuracy_score
    accuracy_score(y_train, y_train_pred)
```

```
0.9108910891089109
```

Fig 7.2.3.accuracy of train data

Print the classification report and check the accuracy of the training data

```
[ ] from sklearn.metrics import classification_report, confusion_matrix
    print(classification_report(y_train, y_train_pred))
```

```

      precision    recall  f1-score   support

     0       0.92      0.88      0.90         92
     1       0.90      0.94      0.92        110

 accuracy          0.91          202
  macro avg       0.91      0.91      0.91          202
 weighted avg     0.91      0.91      0.91          202
```

Fig 7.2.3. classification report of train data

We got an accuracy score of around 91% for the training data which is considered to be a good score.

Now let's check the same for test data.

7.2.4 Predicting on test data

We can predict the testing data using predict function on X_test

```
[ ] y_test_pred = log_reg.predict(X_test)
```

```
[ ] y_test==y_test_pred
```

```
342    True
191    True
349   False
288    True
56     True
...
172    True
391   False
629   False
197    True
720   False
Name: target, Length: 100, dtype: bool
```

Fig 7.2.4. predicting on test data

Import confusion_matrix and classification_report from sklearn.metrics

Visualizing the confusion matrix using heat map




Fig 7.2.4. visualizing confusion matrix of test data

Find the accuracy score

```
[ ] accuracy_score(y_test, y_test_pred)
```

0.81

Fig 7.2.4. accuracy score of test data

Print the classification report and check the accuracy of the training data

```
from sklearn.metrics import classification_report, confusion_matrix
print(classification_report(y_test, y_test_pred))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.79 | 0.80 | 0.80 | 46 |
| 1 | 0.83 | 0.81 | 0.82 | 54 |
| accuracy | | | 0.81 | 100 |
| macro avg | 0.81 | 0.81 | 0.81 | 100 |
| weighted avg | 0.81 | 0.81 | 0.81 | 100 |

Fig 7.2.4. classification report of testing data

We got an accuracy score of around 81%, therefore we can say that it is a best fit and the model predicted very well.

Approach 3:

7.3.1. KNN CLASSIFIER:

A supervised machine learning algorithm is one that relies on labeled input data to learn a function that produces an appropriate output when given new unlabeled data.

The k-nearest neighbors (KNN) algorithm is a simple, easy-to-implement supervised machine learning algorithm that can be used to solve both classification and regression problems. However, it is more widely used in classification problems in the industry. It is

used for classification and regression of known data where usually the target variable is known beforehand.

K nearest neighbors is an algorithm that stores all available cases and classifies new cases based on a similarity measure (e.g., distance functions). It should also be noted that all three distance measures are only valid for continuous variables.

The 'k' stands for the number of nearest neighbors for the newly entered value.

The kNN algorithm assumes that similar things exist in close proximity. In other words, similar things are near to each other. kNN captures the idea of similarity (sometimes called distance, proximity, or closeness) with some mathematics we might have learned in our childhood— calculating the distance between points on a graph.

This works based on minimum distance from the query instance to the training samples to determine the k-nearest neighbors. After we gather these k-nearest neighbors, we take the simple majority of these k nearest neighbors to be the prediction of the query instance.

Methods of calculating distance between points:

The **first step** is to calculate the distance between the new point and each training point. There are various methods for calculating this distance, of which the most commonly known methods are – Euclidean, Manhattan and Hamming distance.

1. **Euclidean Distance:** Euclidean distance is calculated as the square root of the sum of the squared differences between a new point (x) and an existing point (y).
2. **Manhattan Distance/City Block Distance:** This is the distance between real vectors using the sum of their absolute difference.

7.3.2 Train the model

Import KNN and create an object for that class

```
[ ] # Model Building:
    from sklearn.neighbors import KNeighborsClassifier
    knn = KNeighborsClassifier(n_neighbors=40, metric='euclidean')

    # Apply the knn object on the dataset(Training Phase)
    # Syntax: objectName.fit(Input, Output)
    knn.fit(scaled_X_train, y_train)

☞ KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='euclidean',
                       metric_params=None, n_jobs=None, n_neighbors=40, p=2,
                       weights='uniform')
```

Fig 7.3.2. importing KNN

7.3.3 Predicting on Train data

```
[ ] # Predictions on the data
#predict function--> gives the predicted values
# Syntax:objectname.predict(Input)
y_train_pred = knn.predict(scaled_X_train)
y_train_pred

array([1, 1, 1, 1, 1, 1, 0, 1, 0, 0, 1, 0, 1, 0, 0, 1, 1, 1, 1, 1, 1, 1,
       0, 0, 1, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1, 0, 0, 0, 1, 0, 1, 1, 1, 0,
       0, 1, 1, 0, 1, 0, 0, 0, 1, 1, 1, 1, 0, 0, 1, 1, 0, 0, 1, 1, 0, 1,
       1, 0, 1, 1, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 0, 1, 1, 1, 0, 0, 1, 0,
       0, 1, 1, 0, 1, 0, 1, 0, 0, 1, 1, 0, 1, 0, 0, 1, 1, 1, 1, 1, 0, 0,
       0, 0, 0, 1, 1, 1, 0, 0, 1, 0, 0, 1, 1, 1, 1, 0, 0, 1, 0, 1, 1, 1,
       1, 0, 0, 1, 0, 1, 0, 0, 0, 1, 1, 0, 1, 0, 0, 0, 1, 1, 1, 1, 0, 0,
       1, 1, 0, 1, 1, 0, 1, 0, 1, 1, 1, 0, 0, 0, 1, 1, 0, 1, 1, 1, 1, 1,
       1, 1, 0, 1, 0, 1, 1, 0, 0, 1, 1, 1, 1, 1, 1, 1, 0, 1, 0, 1, 0,
       1, 1, 1, 0])
```

Print the classification report and check the accuracy of the data

```
[112] # Check the accuracy, classification report
from sklearn.metrics import classification_report
print(classification_report(y_train, y_train_pred))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.89 | 0.82 | 0.85 | 92 |
| 1 | 0.86 | 0.92 | 0.89 | 110 |
| accuracy | | | 0.87 | 202 |
| macro avg | 0.87 | 0.87 | 0.87 | 202 |
| weighted avg | 0.87 | 0.87 | 0.87 | 202 |

Fig 7.3.4. classification report of data


```

▶ from sklearn.metrics import accuracy_score
# Checking for optimum k-value
# Build the models with multiple k values
scores=[]
for k in range(1, 20):
    knn_model = KNeighborsClassifier(n_neighbors=k)
    knn_model.fit(scaled_X_train, y_train)
    pred_test = knn_model.predict(scaled_X_test)
    scores.append(accuracy_score(y_test, pred_test))
scores

```

```

↳ [0.73,
    0.73,
    0.79,
    0.79,
    0.76,
    0.77,
    0.75,
    0.78,
    0.77,
    0.8,
    0.79,
    0.8,
    ...

```

```
[ ] # Plot of K values and Scores
plt.plot(range(1,20), scores, marker='o', markerfacecolor='r', linestyle='--')
```

[<matplotlib.lines.Line2D at 0x7f78c6966f60>]




Fig 7.3.5. plot a graph

7.3.6 Predicting on Train data

We can predict the training data using predict function on X_train

```
[111] # Prediction on training data
      final_train_pred = final_model.predict(scaled_X_train)
      final_train_pred

↳ array([1, 0, 1, 1, 1, 1, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 1, 1, 1, 1, 1, 1,
        1, 0, 1, 0, 0, 0, 0, 1, 1, 1, 1, 0, 0, 1, 0, 0, 1, 0, 1, 1, 1, 0,
        0, 1, 1, 0, 1, 0, 0, 0, 1, 1, 1, 1, 0, 0, 0, 1, 0, 0, 1, 1, 0, 1,
        1, 0, 1, 1, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 0, 1, 1, 1, 0, 0, 1, 0,
        0, 0, 1, 0, 1, 0, 1, 0, 0, 1, 1, 0, 1, 0, 0, 1, 1, 1, 1, 1, 0, 0,
        0, 0, 0, 1, 1, 1, 0, 0, 1, 0, 0, 1, 1, 1, 1, 0, 0, 1, 0, 1, 1, 1,
        1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 1, 0, 1, 0, 0, 0, 1, 1, 1, 1, 0, 0,
        1, 1, 1, 1, 1, 0, 1, 0, 1, 1, 1, 0, 0, 0, 1, 1, 0, 1, 1, 1, 1, 1,
        1, 1, 0, 1, 0, 1, 1, 0, 0, 1, 1, 0, 1, 1, 1, 1, 1, 0, 1, 0, 1, 0,
        1, 1, 1, 0])
```

Fig 7.3.6 . A.predicting on train data

Import confusion_matrix and classification_report from sklearn.metrics

Visualizing the confusion matrix using heat map

```
[110] #Confusion Matrix of Training data
#Syntax: confusion_matrix(ActualValues, Predicted Values)
from sklearn.metrics import confusion_matrix
sns.heatmap(confusion_matrix(y_train, final_train_pred), annot=True,
            fmt='d', annot_kws={'va':'top','ha':'right'}) # d--> integer formatting
```




Fig 7.3.6.b..visualizing confusion matrix of train data

Find the accuracy score

```
[113] from sklearn.metrics import accuracy_score
accuracy_score(y_train,y_train_pred)
```

0.8712871287128713

Fig 7.3.6.c.accuracy of train data

Print the classification report and check the accuracy of the training data

```
# Classification report for training Data
# Precision--> PPV--> Out of the positive predicted values, how many truely positive
print(classification_report(y_train, final_train_pred))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.91 | 0.88 | 0.90 | 92 |
| 1 | 0.90 | 0.93 | 0.91 | 110 |
| accuracy | | | 0.91 | 202 |
| macro avg | 0.91 | 0.90 | 0.90 | 202 |
| weighted avg | 0.91 | 0.91 | 0.91 | 202 |

Fig 7.3.6. d.classification report of train data

We got an accuracy score of around 91% for the training data which is considered to be a good score.

Now let's check the same for test data.

7.3.6 Predicting on test data

We can predict the testing data using predict function on X_test

```
[104] # Predictions on Test Data
final_test_pred = final_model.predict(scaled_X_test) # y_test
final_test_pred
```

```
array([1, 1, 1, 1, 1, 0, 1, 0, 0, 0, 0, 1, 0, 1, 1, 0, 0, 1, 1, 0, 1, 1,
        1, 1, 0, 0, 1, 1, 0, 1, 0, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1,
        0, 1, 1, 1, 1, 0, 1, 1, 1, 0, 1, 0, 0, 0, 1, 1, 1, 1, 1, 0, 0, 1,
        1, 1, 0, 0, 1, 0, 1, 1, 0, 0, 1, 1, 0, 0, 1, 0, 1, 0, 1, 0, 1, 1,
        1, 0, 1, 1, 1, 0, 0, 1, 1, 1, 1, 1])
```

Fig 7.3.6.a predicting on test data

Import confusion_matrix and classification_report from sklearn.metrics

Visualizing the confusion matrix using heat map

```
[105] # Compare actual values of test data(y_test) and final_test_pred(model predicted values)
      # Confusion_matrix(actualValues, predictedValues)
      sns.heatmap(confusion_matrix(y_test, final_test_pred), annot=True, fmt='d')
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f78bfbabba8>




Fig 7.3.6.b visualizing confusion matrix of test data

Find the accuracy score

```
from sklearn.metrics import accuracy_score
accuracy_score(y_test,y_test_pred)
```

0.81

Fig 7.3.6.c accuracy score of test data

Print the classification report and check the accuracy of the testing data

```
[106] # Classification Report for Test Data  
      print(classification_report(y_test, final_test_pred))
```



| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.78 | 0.63 | 0.70 | 46 |
| 1 | 0.73 | 0.85 | 0.79 | 54 |
| accuracy | | | 0.75 | 100 |
| macro avg | 0.76 | 0.74 | 0.74 | 100 |
| weighted avg | 0.75 | 0.75 | 0.75 | 100 |

Fig 7.3.6 .d classification report of testing data

We got an accuracy score of around 76%,

8. CONCLUSION:

1.logistic regression:81-91%

2.naive bayes classifier:81-88%

3.knn classifier:76-91%

-> In this project, I used Machine Learning to predict whether a person is suffering from a heart disease. After importing the data, I analysed it using plots. Then, I did generated dummy variables for categorical features and scaled other features. I then applied 3 Machine Learning algorithms, K Neighbors Classifier, logistic regression ,naive bayes classifier. I varied parameters across each model to improve their scores. In the end, Logistic Regression achieved the highest score of 81-91% as compared to other algorithms

9. REFERENCES:

- 1.<https://medium.com/code-heroku/introduction-to-exploratory-data-analysis-eda-c0257f888676>
- 2.https://en.wikipedia.org/wiki/Machine_learning
- 3.<https://www.edureka.co/blog/what-is-data-science/>
- 4.<https://www.edureka.co/blog/data-science-applications/>
- 5.https://en.wikipedia.org/wiki/Naive_Bayes_classifier
- 6.https://en.wikipedia.org/wiki/Logistic_regression
- 7.https://en.wikipedia.org/wiki/K-nearest_neighbors_algorithm

8. CONCLUSION:

Comparing the two models we can clearly say that logistic regression performed well with an accuracy of around 93%-95% compared to naive bayes which gave an accuracy score of around 83%-85%.

9. REFERENCES:

1. <https://medium.com/code-heroku/introduction-to-exploratory-data-analysis-eda-c0257f888676>
2. https://en.wikipedia.org/wiki/Machine_learning
3. <https://www.edureka.co/blog/what-is-data-science/>
4. <https://www.edureka.co/blog/data-science-applications/>
5. https://en.wikipedia.org/wiki/Naive_Bayes_classifier
6. https://en.wikipedia.org/wiki/Logistic_regression